

BIDU · FY2026 Q1

2026-05-18

Baidu, Inc. · 28 turns · 13 participants

CAST (IN ORDER OF APPEARANCE)**Operator**

Operator

Robin Li

Co-Founder and CEO

Henry Haigiancu

Chief Financial Officer

Alicia Yap

Analyst, Citigroup

Lincoln Kong

Analyst, Goldman Sachs

Thomas Chong

Analyst, Jefferies

Zhou Shen

EVP, Baidu AI Cloud Group

Juan Lin

Director of Investor Relations

Henry Haigiancu

Chief Financial Officer

Alex Yao

Analyst, JPMorgan

Gary Yu

Analyst, Morgan Stanley

Miranda Zung

Analyst, Bank of America Securities

Wei Xiong

Analyst, UBS

OPERATOR Operator

Hello and thank you for standing by for Baidu's first quarter 2026 earnings conference call. At this time, all participants are in a listen-only mode. After management's prepared remarks, there will be a question and answer session. Today's conference is being recorded. If you have any objections, you may disconnect at this time. I would now like to turn the meeting over to your host for today's conference, Juan Lin, Baidu's Director of Investor Relations.

JUAN LIN Director of Investor Relations

Hello, everyone, and welcome to Baidu's first quarter 2026 earnings conference call. Baidu's earnings release was distributed earlier today, and you can find a copy on our website as well as on Newswire Services. On the call today, we have Rodney Lee, our co-founder and CEO. Julius Rongluo, our EVP in charge of Baidu Mobile Ecosystem Group, MEG. Zhou Shen, our EVP in charge of Baidu AI Cloud Group, ACG. And Henry Haigiancu, our CFO. After our prepared remarks, we will hold a Q&A session. Please note that the discussion today will contain forward-looking statements made under the safe harbor provisions of the U.S. Credit Security Litigation Reform Act of 1995. Forward-looking statements are subject to risks and uncertainties that cause actual results to differ materially from our current expectations. For detailed discussions of these risks and uncertainties, please refer to our latest annual report and other filings with SEC and Hong Kong Stock Exchange. Baidu does not undertake any obligation to update any forward-looking statements except as required under applicable law. Our earnings press release and this call include discussions of certain unaudited non-GAAP financial measures. Our press release contains a reconciliation of the unaudited non-GAAP measures to the unaudited most directly comparable GAAP measures and is available on our IR website at ir.biz.com. As a reminder, this conference is being recorded. In addition, a webcast of this conference call will be available by this IR website. I will now turn the call over to our CEO, Robin.

ROBIN LI Co-Founder and CEO

Generated RMB 26.0 billion in total revenue in Q1 of 2% year-over-year, marking a return to positive growth. Revenue from our core AI-powered business reached RMB 13.6 billion, up 49% year-over-year. For the first time, it

accounted for more than half of Baidu general business revenue, reaching 52%. This is an important milestone as AI-powered business has now become the majority of our revenue mix. AI Cloud Infra delivered exceptional momentum in Q1 with overall revenue growing 79% year-over-year. Within AI Cloud Infra, GPU cloud revenue continued its strong trajectory from last quarter's 143% growth, accelerating further to 184% year-over-year. ApolloGo also had a strong quarter. We delivered 3.2 million fully driverless rides in Q1, sustaining triple-digit growth in total rise year-over-year, reflecting the continued scaling of our operations. Together, these results confirm that AI has clearly become the primary growth driver of Baidu, reinforcing our position as an AI-first company. As AI adoption continues to accelerate, real-world applications are expanding opening up new and increasingly diverse demand for AI capabilities. We are confident in our ability to capture these opportunities as they unfold and believe AI will continue to drive the next phase of Baidu's growth. Now, let me walk you through the key highlights of this quarter, starting with AI Cloud Infra. As AI adoption accelerates across industries, we continue to see demand surge across both training and inference workloads, with inference ramping especially fast and accounting for a growing share of overall demand. Q1 was a quarter of significantly accelerated growth for our AI cloud infra, with revenue growth well above the broader market. The mix of our business continued to shift toward higher quality revenue streams. GPU cloud, which typically carries stronger margins, has become a meaningful contributor to our total AI cloud infra revenue, underscoring the ongoing improvement in overall business health. A key driver behind this momentum is the differentiated advantage of Baidu's full-stack AI capabilities, one that very few companies globally can truly claim. With proprietary components at every layer from underlying infrastructure to applications, we were able to ensure stable and reliable compute supply while also optimizing end-to-end across the entire stack, continuously improving performance, reducing costs, and delivering compelling cost effectiveness for our customers. As AI applications continue to proliferate, this full-stack advantage becomes increasingly pronounced, enabling us to capture a broader and more diverse range of opportunities. At the infrastructure layer, we hold a distinct advantage through Kunlun Xin, our self-developed AI chips. We have seen strong and expanding demand for Kunlun Xin, with a growing number of customers across diverse industries adopting it for a broadening range of AI workloads. This reflects growing market recognition of Kunlunxin's stability, efficiency, compatibility, and versatility. It is also among the first domestic AI chips to achieve large-scale commercial deployment in a single AI computing cluster of over 30,000 accelerators, with industry-leading cluster performance and stability. Built on a comprehensive software stack, Kunlunxin delivers raw compatibility with different models and frameworks, as well as strong usability across enterprise environments. To date, it has been optimized and validated for workloads across various models, covering the latest versions of Ernie and other mainstream foundation models, with inference support recently extended to DeepSeq v4 GLM 5.1 and Minimax M2.7. As an important component of our AI infrastructure, Kunlunxin further strengthens the foundation of our infrastructure layer, enabling Baidu AI Cloud to support customers' AI deployment with greater efficiency, reliability and cost effectiveness, and enhancing the overall competitiveness of our cloud offerings. These advantages are translating into strong client momentum on the infrastructure side. Baidu AI Cloud has become a trusted infrastructure partner for a growing number of major companies across a broad range of industries, including internet, gaming, embodied AI, autonomous driving, smartphones, financial services, and more. This quarter, we added several prominent new clients, including leading model companies. Our client base also includes leading names such as Unitree, Honor, Opal, and Bevo. At the same time, existing top-tier clients continue to deepen their collaboration with us and scale their usage, driving healthy expansion across our client base. On the MOS front, as OpenClaw gained traction across the industry, we moved quickly to expand the model library on our Qianfan MOS platform. In addition to Ernie, Qianfan now supports an expanding set of in-demand models, including popular ones from Zhipu AI, Minimax, Kimmy, and DeepSeek, keeping our model library comprehensive and up-to-date. In March, daily average token consumption from external customers grew to nearly seven times the level of a year ago, while our MOS revenue also scaled rapidly. We believe the MOS platform still has significant untapped potential as the

ecosystem around agents and AI applications continues to evolve. On foundation models, we recently launched Ernie 5.1, which delivers stronger tax capabilities, a more compact model of size, and enhanced the reasoning compared to its predecessor. We also made advances in key areas, including code generation, agentic capabilities, and deep search. Recently on the LM arena, Ernie 5.1 ranks first among Chinese models on the text leaderboard. Ernie 5.1 also topped the LM arena search leaderboard among Chinese models, ranking fourth globally making it the only Chinese model to appear on that leaderboard as well. Looking ahead, we remain firmly committed to advancing Ernie through an application-driven approach, continuously iterating based on real-world needs to keep Ernie at the forefront of AI capabilities. Now let me turn to AI applications. We have long believed that the true value of AI is ultimately realized through applications, and we have been early and persistent in building a comprehensive portfolio serving both enterprises and individual users. This quarter, we continue to see encouraging progress across several high potential directions. Let me highlight a few examples. The first is DoMate, our AI agent for everyday productivity, which we recently showcased at Baidu Create. Do-Mate is designed to execute complex, multi-step workflows across applications and files autonomously, handling long-running tasks from start to finish. Available across both PC and mobile, it enables users to initiate tasks anytime and from anywhere while operating continuously in the background as a 24-7 AI assistant. Users simply describe what they need and come back to results. What truly differentiates DuMate is its seamless integration with Baidu's proprietary skills, including AI Search, Vibe Coding MiaoDa, FaMou, and more. As we continue to expand DuMate's skill ecosystem, we believe it will be able to better tackle an ever-wider range of office workflows and complex, real-world tasks, helping users complete them end to end more effectively. Turning to digital humans, our hyper-realistic digital human technology continued to advance with improved performance and increasing readiness for large-scale deployment. On the cost front, we achieved around 80% cost reduction over the past two quarters, lowering the adoption barrier and making our digital humans more affordable and accessible for a broader range of clients. Meanwhile, we are also taking our digital human capabilities global. At the recent Baidu Create, we launched an overseas digital human platform that enables merchants and creators to easily generate digital human content from e-commerce live streams to digital human videos and beyond. To make our digital humans truly work for global markets, we have built in deep localization from the ground up, supporting 24 languages including Spanish, French, and Thai with scripts and presentation styles culturally adapted to resonate with local audiences. This helps merchants run round-the-clock digital human live streams that feel authentically native, unlocking new levels of efficiency and conversion potential across global markets. Our growing partner base in China and overseas includes Jingdong, Suoyebang, TikTok, and Shopee, with several partners deepening their collaboration with us. Next is MiaoDa, our vibe coding platform. MiaoDa empowers anyone to bring their ideas to life without writing a single line of code, and we are seeing this value increasingly recognized. In March, monthly active users of MiaoDa grew around 70% quarter over quarter, while our domestic paying user rate reached approximately three times the level at the end of last year. At Baidu Create, we launched MiaoDa 3.0, introducing an enterprise version and a mobile app, enabling broader adoption across both individuals and enterprises, as well as more flexible usage across time and use scenarios. Notably, MiaoDa now supports the generation of standalone mobile applications, further expanding what users can create with MiaoDa. Another example is Phamo Agent, our self-evolving agent designed to address complex operational challenges across industries and help enterprises unlock meaningful productivity gains. With the launch of FAMO Agent 2.0 at Baidu Create, we further expanded its accessibility. While earlier versions were primarily used by developers and technical teams, FAMO Agent 2.0 lowers the barrier to entry by enabling domain experts to interact with the agent directly through natural language, no coding expertise required. For example, at Qingdao port, one of the world's leading ports with highly sophisticated scheduling system and deeply complex operational logic, FAMO agent is helping push the efficiency of an already advanced system even further. In an environment where thousands of interdependent variables must be coordinated in real time, FAMO agent autonomously explores the solution space to identify optimal decisions across birth

scheduling, equipment allocation, and cargo prioritization. Even on top of an already highly optimized baseline, FAMO agent continues to unlock incremental efficiency gains, further enhancing the overall operational performance. Next is AI search. In Q1, We continue to advance our AI search transformation with a particular focus on improving user satisfaction and the overall search experience. Through ongoing enhancements in model capabilities, we further improved how search results are planned, structured, and generated, enabling better assessment of content quality, broader distribution of high-quality information, and a significant reduction in low-quality content. Meanwhile, Early Assistant continued to see strong user engagement, driven by ongoing improvements to its interaction experience. In March, daily active users of Early Assistant nearly doubled year over year, while daily average conversation rounds more than tripled over the same period. Next day, retention also improved meaningfully, reflecting stronger user stickiness. Looking ahead, we will continue to deepen the integration between AI Search and Early Assistant, further enhancing the experience across information discovery, content understanding, and task completion. Beyond the digital world, AI is reshaping the physical world in profound ways and Robotaxi stands as the most powerful embodiment of this transformation. In Q1, ApolloGo, our Robotaxi business, maintained strong momentum with fully driverless rides continuing to grow and our safety record remaining industry-leading. We also continued to advance our international expansion, making steady progress across key overseas markets. In Europe, we are on track to commence open road testing in Switzerland, and our first vehicles have arrived in London in preparation for testing with Uber and Lyft, expected to begin soon. In the Middle East, ApolloGo's fully driverless operations are now running across multiple zones in Dubai. Also in late March, we officially launched the ApolloGo app making us the first and only autonomous ride-hailing service with its own standalone app there. Beyond traditional ride-hailing, we are exploring new use cases that broaden ApolloGo's commercial reach. In Hainan, one of China's most popular tourist destinations, we partnered with Car, Inc. on a rental model with our fully driverless vehicles stationed directly at the arrival level of Haiku Airport. So ApolloGo is right there waiting as visitors step out of the terminal. We believe Robotaxi can deliver value well beyond daily commuting, opening up new monetization opportunities in the process. As ApolloGo continues to scale fully driverless operations, we have encountered a broader and increasingly complex range of real-world scenarios, including system and operational complexities that only emerge at larger scale. When such situations arise, we handle them with rigor and use them to continuously strengthen our operations. More broadly, ApolloGo has moved well beyond technology demonstration and small-scale pilots. At this scale, we are addressing a new frontier centered on how robo-taxi services fit more naturally into public transportation, city operations, and everyday life. These experiences are helping us build the expertise and knowledge needed for ApolloGo to coexist more seamlessly with the broader transportation ecosystem over time and ultimately to become a more convenient and trusted service for the people we serve. In closing, our AI-powered businesses delivered strong momentum across the board in Q1. AI has now become the core driver and the majority of our business, and we see this role only growing from here. As we scale AI across an increasingly diversified portfolio and extend our reach into global markets from AI applications to autonomous ride hailing, we see significant opportunities opening up on multiple fronts. We are confident in our ability to capture them. With that, let me turn the call over to Henry to go through the financial results.

HENRY HAIGIANCU Chief Financial Officer

Thank you, Robin, and hello, everyone. In Q1, we continue to make progress on our key priorities, enhancing disclosure, growing our core AI-powered business, and improving operational efficiency. I'd like to highlight a few results from the quarter. Total revenue of Baidu General Business grew 2% year-over-year, returning to positive growth after several quarters of decline. Non-GAAP operating income of Baidu General Business increased 39% quarter-over-quarter to RMB 4.0 billion. Operating cash flow for Baidu remained positive for the third consecutive quarter at RMB 2.7 billion, reflecting the continued improvement in our operating efficiency and overall business health. We also

reached an important milestone. Our core AI-powered business accounted for more than half of Baidu's general business revenue for the first time. Revenue from Baidu core AI-powered business exceeded RMB 13 billion, up 49% year over year. Within this, AI Cloud Infra growth significantly outpaced the broader market, while our AI applications portfolio continued to flourish across multiple fronts. Combined, AI Cloud Infra and AI applications drove our total AI Cloud revenue to RMB 11.3 billion in the first quarter of 2026. Beyond the cloud, ApolloGo further reinforced its position as a global leader in autonomous ride hailing and continued to expand its operations. Collectively, these results point to a business that is becoming both more AI-driven and more financially healthy. Now let me walk through the details of our first quarter 2026 financial results. Total revenue of Baidu was RMB 32.1 billion, decreasing 2% quarter over quarter. Revenue from Baidu general business was RMB 26.0 billion, increasing 2% year over year, and remaining flat quarter over quarter, among which the increase in others was primarily driven by the growth of AI cloud business. Revenue from iQiyi was RMB 6.2 billion, decreasing 8% quarter over quarter. Cost of revenues was RMB 19.6 billion, increasing 7% quarter over quarter, primarily due to an increase in costs related to AI cloud business, partially offset by decreases in content costs and traffic acquisition costs. Operating expenses were RMB 9.3 billion, decreasing 28% quarter over quarter, primarily due to decreases in expected credit losses and the personnel related expenses. Operating income was RMB 3.2 billion and operating margin was 10%. Non-GAAP operating income was RMB 3.8 billion and non-GAAP operating margin was 12%. Total other income net was RMB 626 million compared to RMB 1.2 billion last quarter. Income tax expense was RMB 528 million compared to RMB 1.0 billion last quarter. Net income attributable to Baidu was RMB 3.4 billion. Net margin for Baidu was 11%, and diluted earnings per ADS was RMB 8.76. Non-GAAP net income attributable to Baidu was RMB 4.3 billion. Non-GAAP net margin for Baidu was 14%, and non-GAAP diluted earnings per ADS was RMB 12.06. We defined total cash and investments as cash, cash equivalents, restricted cash, short-term investments, net, long-term time deposits, and how-to maturity investments, and adjusted long-term investments. As of March 31, 2026, total cash and investments were RMB 279.3 billion. Operating cash flow was RMB 2.7 billion. Baidu General Business had approximately 28,000 employees as of March 31st, 2026. With that, operator, let's now open the call to questions.

OPERATOR Operator

Thank you. If you wish to ask a question, please press star 1 on your telephone and wait for your name to be announced. If you wish to cancel your request, please press star 2. If you're on a speakerphone, please pick up the handset to ask your question. Your first question today comes from Alex Yao with JPMorgan. Please go ahead.

ALEX YAO Analyst, JPMorgan

Thank you, Mashman, for taking my question and congrats on the impressive acceleration in AI cloud infra revenue this quarter. Could you guys share more color on the key drivers behind this revenue momentum? And do you have a sufficient compute capacity to support future growth? And then how should we think about the margin profile of AI cloud compared with a more traditional, for example, CPU cloud and the long-term margin trajectory? Thank you.

ZHOU SHEN EVP, Baidu AI Cloud Group

Thank you, Alex. This is Doug. We have seen a remarkably strong enterprise demand for AI infrastructure, both training and inference, with inference showing particularly strong momentum, which is a pretty healthy signal. It tells us that customers have moved beyond training models and are now running AI across more parts of their business, at an accelerating pace. Closely related to this, our MaaS platform is seeing a strong traction. ChenFan is one of the very few MaaS platforms in China. As Roden just mentioned, besides Ernie, we have quickly expanded ChenFan's model library to include the most in-demand models, like Dipsyc, JML, Minimax, and others. And we're seeing a continued growth in token consumption from external customers. Importantly, supporting new models quickly is not a simple plug and play process. It requires high throughput, inference, and efficient model serving capabilities. So we can run these models

reliably at scale and serve more token demand with the same amount of compute. Demand is broad-based across verticals, including air-night, autonomous driving, embodied AI, gaming, advanced manufacturing, and more. It's not just existing customers spending more. We keep winning new ones, too, including industries that historically wouldn't have users of AI and cloud computing, like retail and IP-based consumer brands. addressable market is still expanding and with demand remaining strong and supply relatively tight. We are actively expanding capacity and improving resource efficiency to better support the growing customer needs. Our confidence in capturing this demand comes from our differentiated full-stack AI capabilities which provides two tangible advantages. First, efficiency. Owning and optimizing across the full stack enables us to deliver highly competitive price performance for our customers. And secondly, our proprietary sets have earned strong recognition in real-world deployments. On margins, the key driver is business mix. GPU cloud usually carries better margin profiles than traditional CPU cloud for a few reasons. Firstly, GPU cloud is technically more complex with much higher barriers to entry. Baidu was actually one of the earliest cloud providers in China to build GPU Cloud at scale. And we remain at the forefront. Secondly, demand remains very strong. While high quality supply is relatively tight, customers prioritize proven performance, stability, and availability, not just cost. Certainly, our current AI chips and full stack AI capabilities gives us more room to optimize costs and continued improvement in our customer mix, further supports margin expansion. So if GPU cloud takes a larger and larger share of our total cloud infrastructure revenue, we believe the blending margins for cloud improve structurally. And that's a durable, ongoing trend. So we're confident in the long-term profitability trajectory of our cloud business. Thank you, Alex.

OPERATOR Operator

Your next question comes from Alicia Yap with Citigroup. Please go ahead.

ALICIA YAP Analyst, Citigroup

Thank you. Good evening, management. Thanks for taking my questions. Also, congrats on the solid cloud results. I have a question related to your foundation model. So how does Baidu view the positioning of early models in this increasingly competitive landscape? And looking ahead, what are your investment plans and the key direction for future model iterations? Thank you.

ROBIN LI Co-Founder and CEO

Hi Alyssa, this is Robin. The model landscape is moving very quickly with active releases from players both in China and globally. We believe model capabilities will continue to advance rapidly and strong in-house foundation model capabilities remain essential. So we will continue to invest in earning this conviction. Meanwhile, we have always believed that models ultimately create value through applications. That's why we have consistently taken an application-driven approach. Each iteration of Ernie is guided by real product needs and bin scenarios. Most recently, we released Ernie 5.1, which achieved leading results on LLM arenas, text, and search leaderboards. demonstrating Ernie's continued progress in text capability, reasoning, and search. Going forward, we will continue to iterate Ernie in line with the needs of our key applications, and that's AI Search, Digital Humans, Nelda, and FAMO Agent. These are among the application areas we believe hold the greatest value, and our goal is to build the strongest capabilities where they matter most. For example, we will keep improving Ernie's capability to understand user intent and assess content quality. So AI Search can deliver more accurate, higher quality, and more intelligent results. We also strengthen Ernie's text and multi-model capabilities to make our digital humans more vivid and hyper-realistic, and more effective at engaging users and driving sales in e-commerce live streaming. We enhance coding capabilities to better support live coding, enabling users to build applications through natural language. And as coding becomes more, becomes and increasingly foundational capability in the era. This will be a growing era of focus. And we'll continue strengthen Ernie's capability to identify better and better solutions across complex real-world scenarios,

helping enterprises in a wide range of industries achieve greater efficiency gains. To better support these directions, we have also made organizational adjustments to our model teams and will continue to evolve our structure as needed. We're confident that Ernie will keep getting stronger across all of these areas. Besides Ernie, we also have a range of smaller, faster, and more efficient models, as well as model combinations optimized for specific scenarios. Different use cases have different requirements for capability, cost, latency, and deployment efficiency. Our goal is always to deliver the best outcome for each application. Over the longer term, we believe the full potential of AI applications is still far from realized. As more AI use cases unfold, the value of our application-driven approach will become even clearer, and earnings will become more capable and more valuable along the way. Thank you.

OPERATOR Operator

The next question comes from Wei Xiong with UBS. Please go ahead.

WEI XIONG Analyst, UBS

Sure. Good evening, management. Congrats on very strong cloud momentum, and thanks for taking my question. I want to get your thoughts on the margin side. As AI cloud infrastructure revenue continues to grow rapidly, and now with AI-powered businesses accounting for over 50% of total revenue, how should we think about Baidu's long-term operating margin and the key drivers for margin expansion going forward? Thank you.

HENRY HAIGIANCU Chief Financial Officer

Thank you, Wei. This is Henry.

HENRY HAIGIANCU Chief Financial Officer

In Q1, as you'll notice, our Baidu-called AI-powered business, which mainly includes business beyond traditional online marketing, already exceeded 50% of our total revenue for the first time. This is an important milestone reflecting both AI's growing contribution and a more diversified revenue base. Many of these fast-growing businesses are still scaling and increasing. As they become a larger part of our revenue mix, we expect them to contribute not only to revenue growth, but also to margin expansion, giving us multiple drivers for sustainable profit improvement over time down the road. At this stage, we are investing in the most strategic AI opportunities with conviction. We care a lot about ROI of these investments and believe what we are building today will shape our margin structure for the years to come and create durable competitive advantages. Let me walk through the key businesses where we see this playing out. First of all, as you mentioned, AI cloud infrastructure. The GPO cloud is structurally higher margin than traditional CPU cloud, driven by stronger demand, tighter supply chain, higher technical barriers, and uprisng power. So as it becomes a larger part of our mix, we expect it to be an important driver of margin improvement and expansion. Second, AI applications. This is a naturally high margin business, driven by sticky and subscription-based models, and operating at leverage over time. Third, Robotaxi. Our unit economics have improved consistently since we have achieved a breakeven in Wuhan City. We are still in the investment phase, but the path forward to profitability is becoming clearer as we scale up. In addition, At the corporate level, a few additional levers worth highlighting today. First of all, we continue to drive cost optimization and operational efficiency across entire organizations. And second, we are deploying AI extensively to improve internal productivity as well. So not the least. Third, on the infrastructure side, we are continuously improving server utilization rates which flows directly to the margin over time. So as I mentioned in summary, our revenue mix is rotating towards higher margin faster growing businesses. Our full stack AI capability drives cost efficiencies and the company wide productivity gains compounding over time. We believe the median to long term margin trajectory is compelling for us and we think it's sustainable. Thank you.

OPERATOR Operator

Your next question comes from Gary Yu with Morgan Stanley. Please go ahead.

GARY YU Analyst, Morgan Stanley

Hi. Thank you, management, and congrats again on the strong AI Cloud Infra results. So my question is regarding RoboTaxi. The management provide an update on your overseas RoboTaxi operations, and how should we think about the operating scale as far as the revenue mix between domestic and overseas for RoboTaxi? And how do we think about margin profile comparisons? And in the longer term, how does Baidu see its role in the RoboTaxi ecosystem as an operator, a technology provider, or a platform? Thank you.

ROBIN LI Co-Founder and CEO

Hi, this is Robin. First on scale, ApolloGo remains a global leader. We've completed over 22 million cumulative rides as of April. China is one of the most open markets in the world, so it's natural that our domestic scale today is significantly ahead of overseas markets. But we are also seeing more markets globally opening up to Robotaxi, its regulatory environment turning more positive. We are really happy that our domestic operational experiences prepare us well for international expansion. We have made a significant progress in a very short period. We only began accelerating our international expansion a few quarters ago, and our footprint has expanded across key markets in Europe, Middle East, and Asia. That pace reflects the scalability of both our technology and our operations across different market environments. Our confidence in overseas expansion is backed by the large-scale operational capabilities we've proven out in China. Through years of real-world, fully driverless operations, we have accumulated deep experience in complex road conditions, operational challenges, and corner cases that only emerge at certain scale. And this is not something that can be built overnight. These experiences have continuously sharpened our algorithms and operational standards, making our Robotaxi operation progressively more robust with every mile driven. So when we expand globally, that accumulated experience travels with us and helps us move faster. A good example is our progress from Hong Kong to London. Hong Kong has been an important ride and drive robot taxi market for us. Over the past year plus, we have accumulated valuable experience there and that know-how has helped support our recent entry into London. another major right-hand drive market. Regarding to profitability, ApolloGo has already achieved UE break-even in its largest operational city in China, despite very low fare levels. As we expand globally, the pricing environment becomes much more attractive. We believe our overseas operations have the potential to deliver much stronger profitability as they continue to ramp up. And the overall international market outside of U.S. and China is also bigger than China's domestic market, I think. Finally, our long-term role in the ecosystem, I think it's still early to call. The robotics industry is still evolving. and both the value chain and Vince models are still taking shape. What we focused on right now is continuing to scale, deepening our technology and operational advantages, and maintaining our global leadership. With that foundation in place, we all have the strategic flexibility to define our role as the ecosystem mature and capture long-term value accordingly. Thank you.

OPERATOR Operator

The next question comes from Thomas Chong with Jefferies. Please go ahead.

THOMAS CHONG Analyst, Jefferies

Hi, good evening. Thanks, management, for taking my questions and congratulations on a solid set of results. As AI investment continues to ramp up across the industry, how should we think about Baidu CapEx level in 2026? And how does management prioritize capital allocation between AI investment and shareholders' returns? And finally, could management provide updates about the company's potential for dual primary listing in Hong Kong?

HENRY HAIGIANCU Chief Financial Officer

Thank you. Thank you, Thomas. This is Henry.

HENRY HAIGIANCU Chief Financial Officer

On the topic mentioned on CapEx, our overall approach is to maintain strategic investment intensity while preserving financial discipline. AI remains Baidu's most important long-term opportunity, and we will continue to invest in foundation models to stay competitive at the frontier, but also across our full stack AI capability more broadly. From a financial standpoint, we have the capacity to support this level of investment. Thomas, as you noticed, our operating profit and operating cash flow remain healthy. with our total cash position also at a healthy level. Our operating cash flow for Baidu continues to be positive at 2.7 billion RMB in Q1 this quarter, the third consecutive quarter since turning positive in Q3 last year. And also meanwhile, we are drawing on a mix of financing channels and different instruments, including operating leases, financial leases, and other low-cost bank borrowings as well to fund our AI investments while remaining and maintaining a sound cash position, a healthy balance sheet, and a long-term financing structure. And also, as you mentioned, our shareholder returns, it is also our priority, and we still keep as it is. Last quarter, we announced a new buyback program and also introduce the first dividend policy. We will continue to balance long-term AI investment with shareholder returns, and we remain committed to creating sustainable value for our shareholders. The last point regarding the Hong Kong dual primary listing, as you mentioned, we continuously evaluate initiatives that would help unlock long-term value for the company, including capital market initiatives. We believe Baidu's underlying value is substantial, and we will be flexible and proactive in exploring ways to surface it. We will consider market conditions, regulatory requirements, and shareholder interests, and we will communicate with the market when there are meaningful developments. Thank you.

OPERATOR Operator

Your next question comes from Miranda Zung with Bank of America Securities. Please go ahead.

MIRANDA ZUNG Analyst, Bank of America Securities

Good evening. Thank you for taking my question and congrats on the very strong results. So my question is about the AI agent. So as we are quickly moving into the agentic AI era, how to think about Baidu's strategy for AI applications and agents? What's management's view on the monetization methods? For example, like the token-based pricing or subscription-based model? or project-based model. Thank you.

ROBIN LI Co-Founder and CEO

Hi, Miranda. This is Robin. As I talked about at the Baidu Create conference last week, over the past three years, the biggest AI moments were driven by model breakthroughs. But this year is different. For the first time, it is an agent. an AI application that has captured the world's attention. That shift validates something we've long believed. In the AI era, value ultimately gets realized at the application layer. AI applications have always been a strategic priority for us. and we have built a portfolio around real user needs and bin scenarios, spanning AI-native products and AI transformations of our existing products, serving individuals, enterprises, and industry verticals alike. Within this portfolio, there are several directions we see particularly valuable. The first is search. our largest consumer-facing product, and one we have been consistently transforming with AI. We are pioneering the AI search experience globally, bringing the latest AI capabilities to our hundreds of millions of users and enabling search to deliver more intelligent, structured, and genuinely helpful answers at scale. And the second is digital humans. We continue to push toward higher realism, lower cost, and greater scalability. As costs come down, digital humans are becoming much more accessible to merchants and can be deployed at larger scale. In e-commerce live streaming, they are proving increasingly effective at driving engagement and conversion with performance in many cases comparable to or even

better than human hosts. And third ones, MiaoDa, are a wide coding product. As its capabilities evolve, MiaoDa supports a broader range of applications and workflows, lowering the barrier for users and enterprises to build AI applications through natural language. And fourth is far more agent focused on enterprise scenarios. It helps enterprises navigate complex, dynamic real-world environments, continuously evolving to identify better solutions and drive meaningful efficiency gains. And as it handles more and more complex scenarios across more and more industries, it keeps getting better. So for these four directions, I think the first direction is clearly more advertising-based and the other three directions is more subscription or usage-based or token-based pricing. And I think for search, it's more of a consensus that it's a meaningful market for AI-native applications. And also for the web coding products, they are also comparable ones outside of China. But for digital humans and pharma agents, these are not consensus. These are our own beliefs. We think the Both represent huge market potentials, but not many companies putting into enough resources and efforts to do this kind of thing. And beyond this, we're also exploring new AI-native product forms. And Do-Mate is one recent example. It is a general purpose agent. It brings AI into the workflows and tasks people deal with every day. It is stateful. It remembers a lot of things about you. So it's going to be a quite sticky application, we think. And our monetization The industry is at a very early stage globally and the monetization models are still evolving. Today, token-based pricing is more common. People are essentially paying for foundation models. But over time, AI applications and agents will become more capable of completing real tasks like a human being. We believe monetization will become broader and more results-oriented. That means users are paying for productivity gains, time savings, better experiences, and tangible results. So in the future, people will pay for agents or applications, and the market for this should be much larger than tokens. Thank you.

OPERATOR Operator

Your next question comes from Lincoln Kong with GS. Please go ahead.

LINCOLN KONG Analyst, Goldman Sachs

Thanks, management, for taking my question, and congrats on the solid results. I wonder, could management share your view on the growth outlook and the competitive landscape for the domestic AI chips in China? So how is Kunlun Xing positioned within this market, and what recent demand trends are we seeing, actually? Thank you.

ZHOU SHEN EVP, Baidu AI Cloud Group

Thank you for your question. This is Dou. I think China's domestic homegrown AI chip market is still early, but moving fast. We're seeing a structural shift in AI compute demand, from a chaining heavy to a growing mix of inference. As agentic applications, industry-specific use cases and new forms of applications continue to emerge. Inference is becoming a growing part of the picture. The open client is a good example, driving a wave of inference demand that is higher frequency, more real-time, and more diverse. On supply, domestic AI chips still face near-term challenges around the capacity and the supply chain maturity, partly because demand is growing faster than supply. But over the long term, China's semiconductor industry is developing quickly, supported by a strong manufacturing and supply chain foundation. We believe domestic supply capabilities will keep improving over time. If that happens, competition will increasingly depend on not just on delivering chips, but on whether those chips can perform reliably and efficiently across diverse real-world workloads. Domestic chips are still catching up with the most advanced global products in certain frontier training scenarios, but inference is an area where domestic chips can be highly relevant and competitive. As inference continues to grow, that advantage becomes increasingly relevant. For enterprises, it's not only about the peak chip performance. What matters more is the stability at scale, compatibility with mainstream models and frameworks, migration cost and friction, support for large-scale cluster deployment, and ultimately, cost efficiency. So we think the market will increasingly consolidate around players who can deliver on all of these dimensions. Kunlunxin is well positioned on each of these fronts. Furthermore, Kunlunxin is not just a standalone chip product. It's a critical

part of Baidu's full-stack AI capabilities, spanning from infrastructure to applications. we can continuously optimize across the entire stack, improving model efficiency, reducing inference costs, and delivering AI infrastructure that is more cost-effective, stable, and easier to deploy. We are seeing a strong and growing customer demand for Kunshin, with adoption expanding across industries. So we believe Kunshin is well positioned to capture the opportunities ahead. Thank you.

OPERATOR Operator

Thank you. There are no further questions at this time. And that does conclude our conference for today. Thank you for participating. You may now disconnect.