

DDOG · FY2026 Q1

2026-05-07

Datadog, Inc. · 60 turns · 15 participants

CAST (IN ORDER OF APPEARANCE)**Operator**

Operator

Yuka Broderick

Senior Vice President of Investor Relations

David Obstler

Chief Financial Officer

Analyst

Barclays

Fatima Boulani

Analyst, Citi

Greg Moskowitz

Analyst, Mitsui

Mark Murphy

Analyst, JP Morgan

Peter Weed

Analyst, Bernstein Research

Operator

Operator

Olivier Plamel

Co-founder and Chief Executive Officer

Analyst

Morgan Stanley

Carl Kersted

Analyst, UBS

Gabriella Borges

Analyst, Goldman Sachs

Kurt

Analyst, Evercore

Patrick Colvin

Analyst, Scotiabank

OPERATOR Operator

Good day and thank you for standing by. Welcome to the first quarter 2026 Data Dog Earnings Conference Call. At this time, all participants are in a listen-only mode. After the speaker's presentation, there'll be a question and answer session. To ask a question during the session, you'll need to press star 11 on your telephone. You will then hear an automated message advising your hand is raised. To answer your question, please press star 11 again. Please be advised that today's conference is being recorded. I would now like to turn the conference over to Yuka Broderick, Senior Vice President of Investor Relations. Please go ahead.

YUKA BRODERICK Senior Vice President of Investor Relations

Thank you, Lisa. Good morning, and thank you all for joining us to review Datadog's first quarter 2026 financial results, which we announced in our press release issued this morning. Joining me on the call today are Olivier Plamel, Datadog's co-founder and CEO, and David Obstler, Datadog's CFO. During this call, we will make forward-looking statements, including statements related to our future financial performance, our outlook for the second quarter and the fiscal year 2026, and related notes and assumptions, our product capabilities, and our ability to capitalize on market opportunities. The words anticipate, believe, continue, estimate, expect, intend, will, and similar expressions are intended to identify forward-looking statements and similar indications of future expectations. These statements reflect our views today and are subject to a variety of risks and uncertainties that could cause actual results to differ materially. For discussion of the material risks and other important factors that could affect our actual results, please refer to our Form 10-K for the year ended December 31, 2025. Additional information will be made available in our upcoming Form 10-Q for the fiscal quarter ending March 31, 2026 and other filings for the SEC. This information is also available on the Investor Relations section of our website along with a replay of this call. We will discuss non-GAAP financial measures, which are reconciled to their most directly comparable GAAP financial meas-

ures, in the tables in our earnings release, which is available at investors.datadoghq.com. With that, I'd like to turn the call over to Olivier.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Thanks, Yuka, and thank you all for joining us to go over a very strong start to 2026. Let me begin with this quarter's business drivers. I'm very pleased to say that our teams executed very well and delivered revenue growth of 32% year-over-year, accelerating from 29% last quarter and 25% in the year-over-quarter. We showed broad-based acceleration of revenue growth across cohorts, including both our AI and non-AI customers. Our AI-native customers cohort continued to grow and diversify rapidly, both in the number of customers we serve and the scale of those customers. And this quarter included new land deals with two of the world's biggest AI research teams, helping them improve and optimize their training workflows. I'll talk more about that in a bit. Even more impressive was the growth in our non-AI customers. Non-AI customer revenue growth accelerated again this quarter to meet 20% year-over-year, up from 23% last quarter and 19% in the year-over-year quarter. We think this is a sign of strong continued cloud migration, greater adoption of our products, and customers of all kinds accelerating their use of AI. Finally, churn has remained low, with gross revenue retention stable in the mid to high 90s, highlighting the mission-critical nature of our platform for our customers. Regarding our Q1 financial performance and key metrics, revenue was \$1.01 billion, an increase of 32% year over year and above the high end of our guidance range. We ended Q1 with about 33,200 customers, up from about 30,500 a year ago. We also ended with about 4,550 customers with an ARR of \$100,000 or more, up from about 3,770 a year ago. These customers generated about 90% of our ARR. And we generated free cash flow of \$289 million, with a free cash flow margin of 29%. Turning to product adoption, our platform strategy continues to resonate in the market. For example, products up from 51% a year ago. 35% of our customers use six or more products up from 28% a year ago and 20% of our customers use eight or more products up from 13% a year ago. So we're lending more customers and delivering value across more products and our business continues to grow and Our total ARR now exceeds \$4 billion, and our quarterly revenue exceeded \$1 billion for the first time in Q1. This is a big achievement for all of us at Datadog, and is a product of years of investment in building and innovating for our customers. But we're still just getting started. Of our 26 products, five are over \$100 million in ARR, and another three are between \$50 million and \$100 million in ARR. We're working hard to build and deliver further growth in those products. And this leaves 18 other products which are earlier in their life cycles. We believe each has a potential to grow to more than \$100 million over time. Moving on to R&D, our engineers, enabled with the latest AI coding tools, are building rapidly to help our customers confidently and securely deploy their applications. So let me speak to a few of our product launches this quarter. Let's start with AI. As a reminder, we're talking about our AI efforts in two buckets, AI for Datadog and Datadog for AI. So first, AI for Datadog. These are AI products and capabilities that make the Datadog platform better and more useful for our customers. In March, we launched our MCP server for general availability. With MCP server, Developers access live production data to debug their applications directly in their AI coding agent, or IDE. We delivered BIT AI Security Agent, which autonomously manages Datadog cloud SIM signals, conducts in-depth investigations of potential threats, and delivers actionable recommendations. We've seen BIT AI Security Agent reduce investigations that could take hours to as little as 30 seconds. We also shipped BIT Assistant, now in preview, which helps customers search and act across Datadog using natural language balls. Moving on to Datadog for AI. This includes Datadog capabilities that deliver end-to-end observability and security across the AI stack. We launched GPU monitoring, enabling teams to understand GPU fleet utilization, workload efficiency, thermal and power behavior, and interconnect performance. This drives higher GPU ROI and operational reliability. Our customers continue to move forward with their AI activities, and we can see that in their usage of the Datalog platform. We now have over 6,500 customers sending data for one or more of our AI integrations. Though this is only 20% of total customers, they represent about 80% of our ARR. And our custom-

ers' usage of AI within Datalog platform continues to grow rapidly. Meet AI SRE agent investigations, have more than doubled from December to March. The number of spans sent to our LLM observability product nearly tripled quarter over quarter. The number of Datadog MCP server tool calls quadrupled quarter over quarter. And the number of beef assistant messages increased by a factor of 12 in that period. While we are aggressively building with and for AI, we also continue to expand the Datadog platform to deliver against our customers' increasingly complex needs. to speak to a few of these efforts. Last month, we launched experiments for general availability. Experiments work hand-in-hand with our feature flagging product and combine best-in-class statistical methods with real-time observability guardrails so companies can test for impact, choose among alternatives quickly, and ship with confidence. In addition, our customers now benefit from APM recommendations by analyzing telemetry data from application performance monitoring, resource monitoring, profiler, and database monitoring. APM recommendations automatically identify performance and reliability issues, and most importantly, explain how to fix them. And we announced our plans to launch our next data center in the UK. We see a large opportunity to serve our British customers as cloud adoption accelerates in regulated industries. Last but not least, We are pleased to have received FedRAMP High certification from the U.S. federal government. With this certification, we can now move forward with federal agency customers that require FedRAMP High to handle sensitive workloads. Meanwhile, we continue to expand our product offerings, go-to-market teams, and channel partnerships for public sector customers, both in the U.S. and internationally. So our teams were hard at work again. and we're looking forward to sharing many new products and feature announcements at our Dash User Conference on June 9th and 10th in New York City. Now let's move on to sales and marketing and highlight some of the deals we close this quarter. First, we landed two large deals, a seven-figure and an eight-figure annualized deal with a with the AI research divisions at two of the world's largest technology companies. These organizations are building and training the most advanced AI models in the world. It is critical for them to reduce engineering friction and increase training velocity. But fragmented internal and open source tooling made it harder to identify and solve issues and reduce engineering and research productivity. By using Datadog, both companies are accelerating their pace of innovation on their hyperscale AI training workloads. and this includes optimizing their workflows using GPU monitoring on large parallel GPU grids. Next, we signed a seven-figure annualized expansion for an eight-figure annualized deal with a leading online recruiting platform. This customer is centralizing on Datadog to reduce complexity, drive developer velocity, and improve efficiency. With this expansion, they will replace a standalone tool with Datadog LLM observability to correlate LLM signals with APM and user experience data. This customer will grow to 16 Datadog products, including Datadog MCP server. Next, we signed a seven-figure annualized expansion for an eight-figure annualized deal with a Fortune 500 bank. With this expansion, this customer will migrate their remaining log data into Datadog, fully replacing their legacy log vendor. Most notably, our flex logs give them granular control over cost while meeting strict compliance requirements. This customer uses 10 Datadog products, including Bits.ai SRE agent, to accelerate incident response with AI. Next, we signed a 70-year analyzed expansion with a leading global edge fund. This customer operates thousands of on-prem hosts and network devices. At that scale, their open source monitoring stack has become operationally unsustainable impacting portfolio managers and investment analysts. With this expansion, they will replace their entire on-prem observability layer with Datadog infrastructure monitoring and network device monitoring, and will have unified visibility across their cloud and on-prem environments. This customer will expand to 11 Datadog products. Next, we landed a six-figure annualized deal with a Fortune 500 insurance company. This company's fragmented observability stack led to long outages with incidents reported first by their customers instead of their tooling. By using Datadog and consolidating three legacy APM tools, they expect to move from reactive responses to proactive incident detection. They will adopt 10 Datadog products to start, including all three pillars and LLM observability. Next, we signed a 70-year annualized expansion with one of the world's largest travel groups in APAC. This customer was using Datadog on one business unit But in two others, they were juggling multiple tools and lacked actionable insights. By consolidating

six legacy open source and cloud monitoring tools, the customers save money and improve platform resiliency and performance. This multi-year commitment positions Datadog as their strategic observability provider. And finally, we landed a six-figure annualized deal with a leading Latin American fintech company. This customer serves tens of millions of users across critical financial flows. Their rapid growth outpaced their fragmented front-end monitoring setup, and outages exposed them to financial, operational, and reputational risks. By adopting our digital experience monitoring suite, including RAM, synthetics, and product analytics, they now have full visibility over user activity. With the cost control they also previously lacked. This customer will start with five Datadog products. And that's it for our wins. Congratulations again to our entire go-to-market organization for a great Q1. Before I turn it over to David for a financial review, I want to say a few words on our longer-term outlook. We are pleased with the way we started 2026, as we support our customers' inflection in AI usage and application development, and as they lean into our AI innovations, including Bits AI SRE Agent, Bits AI Security Analyst, Bits Assistant, Datalog MCP Server, GPU monitoring, and many more. There is no change to our overall view that digital transformation and cloud migration are long-term secular growth drivers for our business. But we now have an additional secular growth driver with AI, as we help our customers deliver more value with this transformative new technology. Now more than ever, we feel ideally positioned to help customers of every size and every industry as well as all types of users, whether humans or AI agents, so they can transform, innovate, and drive value through AI and cloud adoption. And with that, I will turn it over to our CFO, David.

DAVID OBSTLER Chief Financial Officer

Thanks, Olivier. This was a very strong quarter for Datadog. Our Q1 revenue was \$1.01 billion, up 32% year over year. Our 6% quarter-over-quarter revenue growth is the highest for Q1 since 2022, and our \$53 million quarter-over-quarter revenue added is the highest ever for Q1. That included the strongest quarter of sequential usage growth from existing customers since the first quarter of 2022. We also delivered an all-time record for sequential ARR added to the quarter. ARR growth accelerated in each month of Q1, and we see a continuation of these healthy growth trends in April. We also achieved strong new logo bookings. New logo annualized bookings a new all-time record by a significant margin and more than double versus a year ago quarter. These included wins in observability and included some of our newer products like security, data observability, and FlexLogs. And our new logo average land size also set a record and more than doubled year over year as we continue to land larger deals. Revenue growth accelerated with our broad base of customers, excluding the AI natives, to mid-20s percent year-over-year, up from 23 percent last quarter and 19 percent in the year-ago quarter. We saw robust growth across our customer base with broad-based strength across customer size, spending bands, and industries. Meanwhile, Our AI-native customer growth continues to significantly outpace the rest of the business. This group continues to diversify and grow, including 22 customers spending more than \$1 million annually and five spending more than \$10 million annually. This group includes the leading companies in foundational models, co-gen tools, and vertical-specific AI solutions. Next, regarding our retention metrics, our trailing 12-month net revenue retention percentage was in the low 120%, up from about 120 last quarter. And our trailing 12-month gross retention percentage remains in the mid to high 90s. Now, moving on to our financial results. Billings were \$1.03 billion, up 37% year over year. and remaining performance obligations, or RPO, was \$3.48 billion, up 51% year-over-year, with current RPO growing in the mid-40s percent year-over-year. RPO duration increased year-over-year as the mix of multi-year deals increased in Q1. As a reminder, we continue to believe revenue is a better indicator of our business trends than Billings and RPO, given their variability. Now, let's review some of the key income statement results. Unless otherwise noted, all metrics are non-GAAP, and we have provided a reconciliation of GAAP to non-GAAP financials in our earnings release. First, Q1 gross profit was \$807 million, with a gross margin of 80.2%. This compares to a gross margin of 81.4% last quarter and 80.3% in the year-ago quarter. As we've discussed in the past, our gross margin varies from quarter to quarter,

with investments into innovations for our customers offset by efficiency efforts. Our Q1 OPEX grew 31% year-over-year versus 29% last quarter and 29% in the year-ago quarter. As a reminder, we continue to grow our investments to pursue our long-term growth opportunities, and this OPEX growth is an indication of our execution of our hiring plans. Q1 operating income was \$223 million for a 22% operating margin compared to 24% last quarter and 22% in the year-ago quarter. Turning to the balance sheet and cash flow statements, we ended the quarter with \$4.8 billion in cash, cash equivalents, and marketable securities. Our cash flow from operations was \$335 million in the quarter. After taking into consideration capital expenditures and capitalized software, free cash flow was \$289 million, and free cash flow margin was 29%. And now for our outlook for the second quarter and for the fiscal year 2026. First, our guidance philosophy overall remains unchanged. As a reminder, we base our guidance on trends observed in recent months and apply conservatism on these growth trends. In addition, as with last quarter, we are applying a higher degree of conservatism to our largest customer. So for the second quarter, we expect revenues to be in the range of \$1.07 to \$1.08 billion, which represents a 29% to 31% year-over-year growth. This guidance implies sequential revenue growth of \$64 to \$74 million, or 6% to 7%, due to the strong growth of revenue in Q1 and into April. Non-GAAP operating income is expected to be in the range of \$225 to \$235 million, which implies an operating margin of 21 to 22%. As a reminder, in Q2, we will be holding our Dash User Conference, which we estimate to cost about \$15 million and which we have reflected in our operating income guidance. Non-GAAP net income per share is expected to be 57 to 59 cents per share based on approximately 369 million weighted average diluted shares outstanding. And for fiscal 2026, we expect revenues to be in the range of 4.3 to 4.34 billion dollars, which represents 25 to 27 percent year-over-year growth. Non-GAAP operating income is expected to be in the range of \$940 to \$980 million, which implies an operating margin of 22% to 23%. And non-GAAP net income per share is expected to be in the range of \$2.36 to \$2.44 per share, based on approximately 372 million weighted average diluted shares outstanding. Finally, some additional notes on the guidance. We expect net interest and other income for fiscal 2026 to be approximately \$170 million. We expect cash taxes for 2026 to be approximately \$30 to \$40 million. We continue to apply a 21% non-GAAP tax rate for 2026 and going forward. And we expect capital expenditures and capitalized software together to be 4% to 5% of revenue in fiscal 2026. To summarize, we are very pleased with our execution in Q1. We are well positioned to help our existing and prospective customers with their cloud migration, digital transformation, and AI adoption journeys. And I want to thank Datadogs worldwide for their efforts. With that, we'll open the call for questions. Operator, let's begin the Q&A. Thanks.

OPERATOR Operator

Thank you. As a reminder, if you would like to ask a question, please press star 11 on your telephone. You'll hear the automated message advising your hand is raised. We also ask that you please wait for your name and company to be announced before proceeding with your question. One moment while we compile the Q&A roster. Our first question today is coming from the line of Mark Murphy of JP Morgan. Please go ahead.

MARK MURPHY Analyst, JP Morgan

Thank you so much and congratulations on an amazing performance. Olivier, is there any way to conceptualize the growth in the sheer raw volume of code that's being produced in the world today due to adoption of code generators such as Cloud Code and Codex and Cursor because they seem to be developing the capability to take on full projects. And some of the charts are showing these capabilities are just exponentially exploding upward in a straight line. I'm wondering how much of that code is going into production and therefore driving activity for Datadog.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Well, we definitely think and see that there's many more applications being created. There's going to be way more complexity in production. We see some of that happening already today. Some of those new applications are getting

into production. They're finding users. We see some signs of that at every layer of our platform. We quoted a few stats on the increasing data volumes we see in our AI products. That's definitely a reflection of that. So we see a... An inflection point there in consumption from customers, we see a move to production that is very real, and we see that across both AI-native and non-AI companies.

MARK MURPHY Analyst, JP Morgan

Okay, and thank you. And as just a quick related follow-up, if we click down one layer, and I'm wondering how you might view the increasing heterogeneity of the environments at the silicon level, because when you look across the Amazon, Vitranium, and Graviton, Google, the TVs, Microsoft has launched the Maya silicon. It looks like that is starting to explode. And, you know, our understanding is that trying to monitor the mixed environments is a lot more difficult than if you just have a uniform fleet of Intel and AMD chips. And we keep hearing all the traditional monitoring tools, they really fail on the custom silicon and Datadog handles it well. And then all this new telemetry, including high bandwidth memory and that type of thing. Can you speak to whether that trend is giving you some tailwind?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Yeah, I mean, look, the broader market that's interesting here is if training, you know, training used to be something only two or three companies were doing or maybe four or five at a large scale. And it looks like training actually might democratize quite a bit more, and many companies will train models on a regular basis. So it becomes more of a viable category for service providers, so a provider like us, basically. I think the heterogeneity of the silicone is definitely a trend that plays in our favor there. You know, the more heterogeneous the more you need someone else to make sense of everything for you and tie it all together and also tie it all with the known GPU aspects and the rest of the infrastructure and the application and the users and the developers, like basically everything we do for it. There's only, you know, when you think of who actually has heterogeneous environments today, that is still a very small number of companies. Google barely just started selling their TPUs to the outside. So I think it's still a small number of companies that are there, but we see a growing opportunity there. Interestingly, last year when we reported earnings, we said we're mostly interested in inference workloads and training is not really a market for us yet. Now we actually see training becoming a market. We started lending customers that are actually hyperscalers, that have a whole host of homegrown technologies, and that are using us specifically in their superintelligence labs to help monitor their workloads, accelerate the training runs, monitor the GPUs also. So we see that as a point of validation that there's going to be a great asset for us.

MARK MURPHY Analyst, JP Morgan

But that's amazing. I think there's a whole new dimension if you can move from – inferencing into the training side, and I caught the reference in the prepared remarks of how you landed a couple of those very large labs, so congrats on everything. Thank you for taking my questions.

OPERATOR Operator

Thank you. One moment for the next question. And our next question will be coming from the line of of Morgan Stanley. Your line is open.

ANALYST Morgan Stanley

Yeah, thank you for taking the question. I want to start it off with David. And this guide to start the year is probably the best we've seen in several years, David. And you laid out the underlying assumptions quite well. I just wanted to do a sanity gut check just on the sort of overall backdrop, macro backup. We do have some geopolitical tensions and those types of things when we think about your mid-East business and any impact from like in your e-commerce or retail

business where there may be some you know, consumer discretionary impacts. I just want to get, like, how you're thinking about those parts of the business, and then I had a follow-up for Ollie.

DAVID OBSTLER Chief Financial Officer

Yeah. We had a very strong quarter across the board. We had, you know, a multi-industry, multi-geography type of quarter, and SMB was very strong. And, you know, the source of our guidance and our raises are, at the core, that type of performance. We haven't seen any particular effect in the consumer businesses or e-commerce businesses yet. We basically have a continuation of trends in those businesses, travels and things like that, that are very similar to the other industries. So we haven't seen it yet. We obviously watch it and look at analytics, but we haven't seen it. In terms of our overall guidance, You know, the trends that we have in organic, we discount across the board, and I think we mentioned our particular treatment of our largest customer.

ANALYST Morgan Stanley

That's very clear. And then, Olivier, for you, I think when we talk to investors about the debate in this category longer term, it's just what does the category look like when agents are doing the triaging, investigating versus human engineers and human SREs? And so what is your sort of vision of how that evolves for Datadog? both from a product standpoint and an experience standpoint from a UI perspective, but also is there going to be new modalities in terms of pricing when agents are consuming the data platform to a higher degree than engineers do today?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Yeah. Look, I think one thing I'd say is it's hard to tell where we're going to be in four or five years. If you had told me two years ago that Most engineers would go back to coding in the console. I would not believe you, and yet, you know, that's one of the winning modalities today. Look, as far as we're concerned, we don't care whether most of the usage is humans, most of the usage is agents, or business model lends itself to it pretty well, like we're usage-based, and it doesn't really matter where the usage is coming from from that perspective. The way we see trends coming up right now is We see both a stratospheric increase of agent usage. So we have a ton of usage on our MCP server. We see customers trying to automate a lot with their own agents using our agents using a combination of those. But we also see an increase of usage of the web interfaces by humans. So right now the two work hand in hand and we keep developing and pushing on both fronts.

OPERATOR Operator

Appreciate that. Thank you. Thank you. One moment for the next question. Next question is coming from the line of Ramo and Chow of Barclays. Please go ahead.

ANALYST Barclays

Thanks and congrats for me as well. One for Olivier and for David. Olivier, if I listen to you and to your prepared remarks, there's a lot of consolidation that people try to do open source tooling. and then realized they kind of needed to come to you and come back. On the other hand, in the industry, we still have a lot of noise around that level. How do you see it in real life? To me, it seems a little bit like observability is just very hot, and then there's different categories where you use certain vendors and some open source tooling. Can you speak what you see in real life there? Thank you.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

I mean, in real life, most companies have open source tools, in some capacity somewhere. But when it comes to having a platform that unifies everything, takes care of everything, does more of the problem-solving for you, that's

typically what customers use us. And the motion we see pretty much everywhere is customers have 4, 6, 7, 15, 25 different things and different pockets in the organization and different business units, and it's a huge mess. and they come to us so they can unify all that. They get better results because all of the data is in one place, the workflows can be automated from end to end, you can get end to end visibility, you don't have blind spots, and also they save money because they don't have all these pockets of inefficiency everywhere. So it's a win for everyone. The thing that's also interesting, in particular this quarter, is that we also landed some large parts of hyperscalers And hyperscalers typically have a culture of building everything themselves and they suddenly have the balance sheet and the human capital to support some of that build up. Like if there was ever a set of companies for whom it makes sense to do it themselves, that would be those companies. And yet, we see that they have the same issues. When it comes to going as fast as they can, being as efficient as they can with their resources, that they come to us to replace some of the things that we were using before.

DAVID OBSTLER Chief Financial Officer

Two things, two metrics to look at that to make the points, Ali, you're making. If you look at our platform adoption and you see both the growth of the different categories and the extension of the categories out to lots of products, that shows you that the consolidation on the Datadog platform you know, has continued and is a very strong trend. And part of that is the movement of solutions, as Ali mentioned, that are both open source, but also the competitive point solutions onto the platform. That's been a significant driver of the revenue growth for some time now, and that continued certainly in Q1.

ANALYST Barclays

Okay, perfect. Thank you. And then, David, for you, last year we did a lot of investments around go-to-market, especially in sales capacity. If you think about now the non-AI category doing better, how much of that is people doing the cloud migrations again, so that's an industry trend, and how much of that is you guys actually being brought a position? Thank you.

DAVID OBSTLER Chief Financial Officer

Yeah, it's a number of things, including One is the expansion of the platform, the consolidation, the successful ramping of sales capacity while not jeopardizing productivity, which has resulted in increasing ARR, and a good environment as well. I think that's what we said last time. There are a number of factors, and certainly what we're proving out here is the investments we've made in go-to-market and are continuing, are paying off, and were the right decision. Ali, anything there?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Yeah, and look, at the end of the day, there's clearly some market tailwinds with the adoption of AI, but also we are outperforming all of our competitors at scale, and we're taking share, and that relates to the structure of our platform, to where we expand with new products, the way these products are maturing and starting to win in their respective categories and the way we've successfully grown the sales capacity.

DAVID OBSTLER Chief Financial Officer

Certainly, the AI investment trend has helped, but what we're trying to do is separate that. And AI investment is probably helping the overall as well, but when you really take that out, you see a very pronounced acceleration here, and that has to do with the factors that I mentioned and Ali talked about.

ANALYST Barclays

Perfect. Congrats. Sounds exciting. Thanks.

OPERATOR Operator

One moment, please, for the next question. Our next question is coming from the line of Gabriella Borges of Goldman Sachs. Please go ahead.

GABRIELLA BORGES Analyst, Goldman Sachs

Hi, good morning. Thank you. Olivia, I find your comments on training versus inference so interesting. Maybe just crystallize for us. Why do you think the training opportunity is happening now or inflecting now and then either for yourself or David? How do we think about the attach rate on training versus inference of observability? If there is a way to benchmark observability spend as a percentage of inference spend, does that number change given the new data that you're seeing on the training side as well? Thank you so much.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

So on the training side, training was very new a couple of years ago. It was something that was only done by very few companies, and it was in a way, very artisanal. It was not a production workload. It was something that researchers were building and that was very one-off and homegrown in ways. And now it's turning into production. It's turning into something that many more companies are doing. It's scaling by orders of magnitude, and it's becoming something that has to be on all the time, reliable, and every minute you lose is a, or rather, every failure you have in your training run is a week you give away to the competition. And so as a result, it becomes way more interesting as a market for Copenhagen. And we see some signs of that. Again, we didn't have a lot of it. We didn't see a lot of it last year. Now, all of a sudden, we're starting to see quite a bit of activity there and demand, and we have success landing with large customers with those products.

DAVID OBSTLER Chief Financial Officer

Yeah, I think going back to the metrics that Ali talked about, you know, in terms of Attach, we said that 6,500 customers are using our integrations, and that's 20% of the customers and 80% of the ARR. So there is Attach. I think it's earlier days for the training. You know, that looks like it will be a contributor, but I think... That's early, and I would sort of look at the larger attachment at this point as the evidence of inference, but also some training.

GABRIELLA BORGES Analyst, Goldman Sachs

Thank you both, and congratulations.

OPERATOR Operator

Thank you. Thank you. One moment, please, for the next question. Our next question is coming from the line of Carl Kersted of UBS. Please go ahead.

CARL KERSTED Analyst, UBS

Okay, great. I wanted to start, Olivier and David, and you can congratulate all of you and the team on reaching that billion-dollar milestone. Well done. David, maybe the question is for you and to hone in specifically on the 2Q guide. Even if you put up a modest beat on that guide, it's going to be, by order of magnitude, the largest sequential dollar ad, I think, in the company's history. And I just wanted to unpack what's giving you that confidence. And in particular, is there anything interesting to call out, David, in terms of the ramp of a couple of the larger research labs, one of which renewed with you guys in the fourth quarter, another one just landed? I presume they're ramping nicely in 2Q, but would love any color. Thank you.

DAVID OBSTLER Chief Financial Officer

Yeah, let me unpack this in a couple ways. As you know, we're a recurring revenue model. So the biggest indication of in the near term of the next quarter is the ARR growth in the previous quarter. And we said we had a record. So essentially at the bedrock of this is sort of the run forward of ARR that we've already signed. The ARR ad was very broad based and was not very concentrated. So, whereas we pointed out some very significant ads, I would say that the first quarter and that AR ad was really diversified and from lots of different places. So, and I think Ali will come in here, but the confidence that we have is you're right. We essentially take what we already have. We discount the growth trends that we've seen. And that produces what you exactly said, which is, you know, whatever your assumptions are on B, you know, a very impressive sequential, really due to what happened in Q1 and the rate of business accumulation by Datadog. Ali, do you want to add?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Yeah, I mean, I basically want to develop on what David just said. The ads were broad-based. I mean, look, and when you look at why we have a great Q1, we also have great customers in Q4. We had talked about it a quarter ago. But even if you take out the customer we landed in Q4 that added the most revenue in Q1, we still had a record quarter in terms of AR ads. So this is really broad-based. And we landed a few more customers in Q1 that don't contribute any revenue yet, but we expect to be big contributors in the future. So when you put all that together, we feel very confident about Q2, hence the numbers you've seen. Yeah, thank you.

OPERATOR Operator

Thank you. One moment, please. And our next question will be coming from the line of Fatima Boulani of Citi. Please go ahead.

FATIMA BOULANI Analyst, Citi

Oh, good morning. Thank you for taking my questions. Ali, I wanted to double back on a question that was asked earlier with respect to telemetry volumes, you know, essentially going parabolic and you are accessing brand new demand vectors, you know, in the foray into training and monitoring, observing training model environments inside some of the world's largest frontier labs. And so I wanted to ask you about the structural changes to the capital intensity of the business. I mean, your CapEx levels are still pretty respectable and pretty muted. So I wanted to get a better understanding of what sort of extrinsic or intrinsic engineering efforts you're undertaking to keep a very efficient CapEx envelope in spite of the fact that it seems like that would increase because of the torrent of telemetry you're seeing on the platform. And then as a related matter, we've seen a rise of sovereign data and data residency requirements kind of ramp as AI models, you know, move into the territory of national security and things like that. So just wondering if you can kind of talk to some of the engineering horsepower internally that you're leveraging to be able to keep a really tight command on capital intensity and, frankly, your gross margins. Thank you.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Yeah, I mean, look, the sort of the investments we're making right now, we run most of our workloads on clouds, meaning you'll see all of that in OPEX, not in CAPEX. So we have low CAPEX. If it changes, we'll tell you, like if for some reason we decide to make different kinds of investments and some of it more out front, some of it more CAPEX, we'll tell you, but that's not the case today. We are definitely ramping up our investments in particular in R&D. and in the scale of the models, we train ourselves on things like that. Right now, there's nothing that you can actually see in the numbers that moves any needle, but if that changes, also we'll tell you. We don't expect any change to our model. So that's on the CapEx side. We're very different businesses in that way from the AI side. On the subject of... data residency and sovereignty of AI, things like that. We definitely see more push for that, more

demand for that in the customer base. And for us, that means investment in two areas. One is in deploying into more geographies and having more certifications to sell to the public sector and to the highest level of the public sector. So we mentioned today data center in the UK, for example, and our federal high certification. We're not stopping there in terms of the certification we're going after with the federal government. So that's an area of investment. Another area of investment is our bring your own cloud products and where we can actually run on our customers' infrastructure. And so we announced that. We released some products there, and we have heavy investment in that area so we can support customers that want to operate in a slightly separate way from the rest of our customers.

OPERATOR Operator

Thank you. Thank you. One moment for the next question.

OPERATOR Operator

And our next question is coming from the line of Kurt of Evercore. Please go ahead.

KURT Analyst, Evercore

Yeah, thanks very much for taking the question, and congrats on a nice start. Ali, I was wondering if you could just give some thoughts on the idea of sort of security for agents. I think one of the big issues in terms of getting agents into production is sort of the security aspect of that. And how do you see Datadog plugging into that opportunity? And then just a quick one for David. Congrats on the FedRAMP reaching that milestone. Are your partner relationships in place to take advantage of this? I realize it will be a long-term opportunity, but just kind of curious how well established you are down there to start seeing some maybe bookings in that area. Thanks, guys.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Yeah, so on the security of agents, we interface with that in two ways. So first, there's the agents we build ourselves because we are building a lot of automation inside of our product for our customers and agents that automatically – identify but also resolve issues without you having to do anything. And a lot of it has to do with understanding what permissions to apply, what kind of guardrails to apply, how to interface with the humans, and how to make that trustworthy and visible in the right way. And so that's pretty much the whole product surface is to doing data. The automation itself actually kind of works already. So you should expect to hear more about that at our conference. This is definitely one big area of investment for us. The security aspects of our engagement. Look, our belief in security is that you need to integrate. You can't just have point solutions that look at one sliver of the whole security posture. You need to get everything all together. And that's one of the areas that we are also covering with our security efforts. So that's part of the whole platform, actually.

DAVID OBSTLER Chief Financial Officer

on the FedRAMP. We've been working on both the different certifications, but at the same time, we've been investing in the go-to-market function, both in terms of reps and channel partners for a number of years. Certainly, there's more investment to be done, but we invested ahead of the certifications because, you know, in this sector, building pipeline, et cetera, takes time. And certainly the channel partner relationships are a very important part of this, and we have been investing but also have more investment to do.

OPERATOR Operator

Thank you.

OPERATOR Operator

One moment, please, for the next question. And our next question is coming from the line of Patrick Colvin of Scotiabank. Please go ahead.

PATRICK COLVIN Analyst, Scotiabank

Thank you, Keith, for taking my question and echoing the congrats of my peers. I guess, you know, Olivia and David, you guys are very deliberate in your messaging on the prepared remarks. And I just want to double check the kind of wording of one of the comments. I think, David, you said a higher degree of conservatism to the largest customer. I guess, did I hear that right? And then, does the higher degree of conservatism reference versus the other customer cohorts, or does it reference versus your guidance philosophy in prior quarters vis-a-vis this customer?

DAVID OBSTLER Chief Financial Officer

It's both. It's the same guidance we used, and we're being very explicit. For All the business except for this largest customers, we've always taken the drivers and discounted them. For this particular customer, we took a higher degree of conservatism than the other part of the customer base and discounted it more. And we were, I think, in the remarks, and you interpreted correct, very explicit, and you're correct.

OLIVIER PLAMEL Co-founder and Chief Executive Officer

I wouldn't give that much weight to the very specific way. We're deliberate, but not all that deliberate. Similarly, both David and I have a raspy voice today, but there's no even meaning to it.

DAVID OBSTLER Chief Financial Officer

But I will remind everybody that we did not change. So if the question also, I think you asked, is did we change or is this a different methodology of both the overall and the large customer than the guidance the last quarter or the previous? The answer is no. It's the same methodology that we've had. So no change, but this has been what we've always been doing.

PATRICK COLVIN Analyst, Scotiabank

Okay. And Olivier, can I ask about your comments about the hyperscalers? Because I thought that was particularly interesting. And the reason why is I don't think you called them out previously before. And they are so prevalent in the modern tech stack today. To your point, they could do this themselves. So I guess how are they using Datadog? Is it for more kind of traditional observability or is it for these newer areas like GPU monitoring that Datadog has performed so well of late?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Well, it's both actually. When you look in general at the large AI customers, they use Datadog the way other companies are largely. with a fairly broad set of our products to cover the full surface of O80T. What's new is we now have a product for GPU monitoring. It's a very new product. And we see the hyperscalers that are coming to us for training workloads in particular being very interested in that. So again, it's too early in the product lifecycle and in the customer lifecycle for these specific customers to call definitive victory there, but we see that as a very encouraging sign of where the market might go in the future, because we think this might be a bellwether of what the next 10, 100, 500 companies that are going to start training workloads are going to want to do. We have some signs that go beyond the customers we signed this quarter that point that way too. Thank you so much.

OPERATOR Operator

Thank you. One moment for the next question, please. And our next question is coming from the line of Peter Weed of Bernstein Research. Please go ahead.

PETER WEED Analyst, Bernstein Research

Thank you. And I'll echo others on the momentum. Great to see. You know, one of, I think, the great successes you talked about was landing a couple of the AI labs for the hyperscalers. Although I think, you know, on the other hand, you've talked in the past around, you know, hyperscalers are typically building observability in-house. What is it really about the AI workloads that are making it more attractive for them to use Datadog? And what might give you confidence that Datadog might be more persistent with them in these types of workloads and as kind of a signal for maybe how other customers might use Datadog around AI, differentiated from things that they might be able to bring in-house other places?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Well, you know, the same reason all of our customers use us, you know, it's a high-stakes, high-complexity and not core. Like, you know, they have to be always differentiated. They can't afford to be late. And it's a really hard job to do all of that. So that's what we build a whole business on. And it's also very true at the highest level for the largest companies.

OPERATOR Operator

Thank you. One moment.

OPERATOR Operator

I'm sorry.

PETER WEED Analyst, Bernstein Research

Yeah, I was just going to say, but I guess the point is you've emphasized that those largest customers have been able to go in-house on some other things. Is there something unique about AI that prevents them from doing that here?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Well, I think the urgency of their development efforts focuses the minds. That's what I would put it. I would say it forces you to figure out what's core and what's not core and what you need to do to maximize your chances of success. And again, it is the same thinking all of our customers have all the time. I think the equation for hyperscalers has often been fairly different because they have, let's call it, unlimited access to staffing. And, you know, they could sort of set their own time horizons for the developments they wanted to make in the situations that would be different with the IRS, maybe. Thank you.

OPERATOR Operator

The line of Greg Moskowitz of Mitsutu. Please go ahead.

GREG MOSKOWITZ Analyst, Mitsutu

All right, great. Thank you, and I'll add my congratulations on a terrific quarter. Just one for me. Holly, I know it's not GA yet, but curious if you have any early feedback on your new cloud-prem offering. As you noted earlier, you know, providing the ability for Datadog to run on customer infrastructure. Could this be another, yet another, I should say, incremental growth opportunity for Datadog? What are your expectations for this?

OLIVIER PLAMEL Co-founder and Chief Executive Officer

Well, definitely, we think, you know, I think there was a question earlier on data residency and, you know, living in a customer's environment. We definitely see a great opportunity there. You know, it is a... a chance that a good portion of the market leans this way in the future. Today, it's not the largest part of the market, but we definitely see the

potential for that. So we're investing heavily in that set of our products. We're starting to see some interest in customer traction there. So we think this can be another growth lever definitely. We also think that it can help us getting into some extremely large-scale workload where customers would not have considered a SaaS offering before, where we can be in the running. So that's very exciting. Great. Thank you. All right. And I think that was our last question. So I want to thank you all for attending the call. And I'll remind you that we have a conference in just a bit more than a month, and I hope to see many of you there. So thank you all.

OPERATOR Operator

This concludes today's program. You may all disconnect.